Claude for Enterprise: Executive Brief

Oct 13, 2026 · @Sean

Bottom line

We recommend a 90-day enterprise pilot of Claude, scored against hard ROI metrics and run under contract terms that keep our proprietary data ours. Claude's edge is rigor: it produces analyses, models and documents accurate enough for a cancer diagnostics business, in a fraction of the usual time.

Why Claude

Claude does best on high-stakes work, where accuracy matters more than speed alone.

Capability	What it does for us
Stress-tests arguments	Challenges strategy, scientific and payer arguments over several rounds before they leave the building
Builds complex work products	Working spreadsheets, financial models, analyses, charts and decks at near-final quality, in hours instead of days
Flags uncertainty	Says when it is unsure instead of inventing an answer, which matters for clinical and regulatory work
Reads large documents	Works across whole protocols, validation packages and sets of papers in one pass
Writes and reviews code	Claude Code builds and reviews bioinformatics and production software
Fits our systems	Connects to GitHub, Google Workspace, email and our data sources

Proving ROI

We prove ROI by measuring the same workflows before and after rollout. We don't rely on anecdotes.

1. **Baseline:** record current hours, cycle time and rework for 4–6 workflows, such as reimbursement models, validation reports, literature reviews, regulatory drafts and board decks.
2. **Pilot:** a 90-day pilot with one defined team on those same workflows.
3. **Measure:** track the metrics below every month.
4. **Report:** a quarterly scorecard to the C-suite, with numbers plus concrete examples.

Metric	How we measure it
Time saved	Hours per deliverable, before vs. after
Cycle time	Days from request to approved output
Quality	Rework rate, errors caught in review, audit findings
Throughput	Deliverables per person per month
Strategic value	Time to submission, grant, payer dossier or launch

```
\text{ROI} = \frac{(\text{hours saved} \times \text{loaded labor cost}) +
\text{value of faster revenue or approvals}}{\text{license cost} +
\text{implementation cost}}
```

Labor savings alone typically cover the license cost. Faster regulatory and commercial timelines are the larger upside.

Protecting our IP

Our proprietary information stays ours through three layers: contract terms, technical controls and our own policies.

Layer	Protection
Contract	Under Anthropic's commercial terms, our data is not used for model training by default

Layer	Protection
Contract	We own the outputs Claude produces for us
Contract	Confidentiality obligations and a data processing agreement
Technical	Single sign-on, automated user provisioning, role-based access and audit logs
Technical	Custom data retention, with zero data retention available for eligible API use
Technical	Independent certifications: SOC 2 Type II, ISO 27001 and ISO 42001
Technical	A BAA (HIPAA business associate agreement) for eligible configurations, required before any patient health information is used
Our policy	Enterprise account only; no personal AI accounts for company data
Our policy	Classify data as allowed, approval required, or prohibited
Our policy	Qualified staff review all clinical, regulatory and external outputs

Legal and IT will confirm these terms against the current Anthropic agreement before signing.

Decision requested

We ask the C-suite to approve a 90-day enterprise pilot. The 90-day results will decide whether we roll out company-wide.

- [ ] Legal and IT review the Anthropic enterprise agreement and the BAA
- [ ] Choose the pilot team and 4–6 baseline workflows
- [ ] Publish the data-use policy before access is granted
- [ ] Review the scorecard at day 90 and decide on rollout